

Linear MIM-Width of Trees [★]

Svein Høgemo, Jan Arne Telle, and Erlend Raa Vågset

Department of Informatics, University of Bergen, Norway.
{svein.hogemo, jan.arne.telle, erlend.vagset}@uib.no

Abstract. We provide an $O(n \log n)$ algorithm computing the linear maximum induced matching width of a tree and an optimal layout.

1 Introduction

The study of structural graph width parameters like tree-width, clique-width and rank-width has been ongoing for a long time, and their algorithmic use has been steadily increasing [11, 17]. The maximum induced matching width, denoted MIM-width, and the linear variant LMIM-width, are graph parameters having very strong modelling power introduced by Vatshelle in 2012 [20]. The LMIM-width parameter asks for a linear layout of vertices such that the bipartite graph induced by edges crossing any vertex cut has a maximum induced matching of bounded size. Belmonte and Vatshelle [2] ¹ showed that INTERVAL graphs, BI-INTERVAL graphs, CONVEX graphs and PERMUTATION graphs, where clique-width can be proportional to the square root of the number of vertices [10], all have LMIM-width 1 and an optimal layout can be found in polynomial time.

Since many well-known classes of graphs have bounded MIM-width or LMIM-width, algorithms that run in XP time in these parameters will yield polynomial-time algorithms on several interesting graph classes at once. Such algorithms have been developed for many problems: by Bui-Xuan et al [4] for the class of LCVS-VP - Locally Checkable Vertex Subset and Vertex Partitioning - problems, by Jaffke et al for non-local problems like FEEDBACK VERTEX SET [14, 13] and also for GENERALIZED DISTANCE DOMINATION [12], by Golovach et al [9] for output-polynomial ENUMERATION OF MINIMAL DOMINATING sets, by Bergougnoux and Kanté [3] for several Connectivity problems and by Galby et al for SEMITOTAL DOMINATION [8]. These results give a common explanation for many classical results in the field of algorithms on special graph classes and extends them to the field of parameterized complexity.

Note that very low MIM-width or LMIM-width still allows quite complex cuts compared to similarly defined graph parameters. For example, carving-width 1 allows just a single edge, maximum matching-width 1 a star graph, and rank-width 1 a complete bipartite graph. In contrast, LMIM-width 1 allows any cut

[★] This is the appendix of our WG submission, the long version with extra figures and full proofs

¹ In [2], results are stated in terms of d -neighborhood equivalence, but in the proof, they actually gave a bound on LMIM-width.

where the neighborhoods of the vertices in a color class can be ordered linearly w.r.t. inclusion. In fact, it is an open problem whether the class of graphs having LMIM-width 1 can be recognized in polynomial-time or if this is NP-complete. Sæther et al [18] showed that computing the exact MIM-width and LMIM-width of general graphs is W-hard and not in APX unless NP=ZPP, while Yamazaki [21] shows that under the small set expansion hypothesis it is not in APX unless P=NP. The only graph classes where we know an exact polynomial-time algorithm computing LMIM-width are the above-mentioned classes INTERVAL, BI-INTERVAL, CONVEX and PERMUTATION that all have structured neighborhoods implying LMIM-width 1 [2]. Belmonte and Vatshelle also gave polynomial-time algorithms showing that CIRCULAR ARC and CIRCULAR PERMUTATION graphs have LMIM-width at most 2, while DILWORTH k and k -TRAPEZOID have LMIM-width at most k [2]. Recently, Fomin et al [7] showed that LMIM-width for the very general class of H -GRAPHS is bounded by $2|E(H)|$, and that a layout can be found in polynomial time if given an H -representation of the input graph. However, none of these results compute the exact LMIM-width. On the negative side, Mengel [15] has shown that STRONGLY CHORDAL SPLIT graphs, CO-COMPARABILITY graphs and CIRCLE graphs all can have MIM-width, and LMIM-width, linear in the number of vertices.

Just as LMIM-width can be seen as the linear variant of MIM-width, path-width can be seen as the linear variant of tree-width. Linear variants of other well-known parameters like clique-width and rank-width have also been studied. Arguably, the linear variant of MIM-width commands a more noteworthy position, since for almost all well-known graph classes where the original parameter (MIM-width) is bounded but other parameters (like clique-width) are not bounded, then also the linear variant (LMIM-width) is bounded.

In this paper we give an $O(n \log n)$ algorithm computing the LMIM-width of an n -node tree. This is the first graph class of LMIM-width larger than 1 having a polynomial-time algorithm computing LMIM-width and thus constitutes an important step towards a better understanding of this parameter. The path-width of trees was first studied in the early 1990s by Möhring [16], with Ellis et al [6] giving an $O(n \log n)$ algorithm computing an optimal path-decomposition, and Skodinis [19] an $O(n)$ algorithm. In 2013 Adler and Kanté [1] gave linear-time algorithms computing the linear rank-width of trees and also the linear clique-width of trees, by reduction to the path-width algorithm. Even though LMIM-width is very different from path-width, the basic framework of our algorithm is similar to the path-width algorithm in [6].

In Section 2 we give some standard definitions and prove the Path Layout Lemma, that if a tree T has a path P such that all components of $T \setminus N[P]$ have LMIM-width at most k then T itself has a linear layout with LMIM-width at most $k+1$. We use this to prove a classification theorem stating that a tree T has LMIM-width at least $k+1$ if and only if there is a node v such that after rooting T in v , at least three children of v *themselves have at least one child* whose rooted subtree has LMIM-width at least k . From this it follows that the LMIM-width of an n -node tree is no more than $\log n$. Our $O(n \log n)$ algorithm computing

LMIM-width of a tree T picks an arbitrary root r and proceeds bottom-up on the rooted tree T_r . In Section 3 we show how to assign labels to the rooted subtrees encountered in this process giving their LMIM-width. However, as with the algorithm computing pathwidth of a tree, the label is sometimes complex, consisting of LMIM-width of a sequence of subgraphs, of decreasing LMIM-width, that are not themselves full rooted subtrees. Proposition 1 is an 8-way case analysis giving a subroutine used to update the label at a node given the labels at all children. In Section 4 we give our bottom-up algorithm, which will make calls to the subroutine underlying Proposition 1 in order to compute the complex labels and the LMIM-width. Finally, we use all the computed labels to lay out the tree in an optimal manner.

2 Classifying LMIM-width of Trees

We use standard graph theoretic notation, see e.g. [5]. For a graph $G = (V, E)$ and subset of its nodes $S \subseteq V$ we denote by $N(S)$ the set of neighbors of nodes in S , by $N[S] = S \cup N(S)$ its closed neighborhood, and by $G[S]$ the graph induced by S . For a bipartite graph G we denote by $\text{MIM}(G)$, or simply MIM if the graph is understood, the size of its Maximum Induced Matching, the largest number of edges whose endpoints induce a matching. Let σ be the linear order corresponding to the enumeration v_1, \dots, v_n of the nodes of G , this will also be called a linear layout of G . For any index $1 \leq i < n$ we have a cut of σ that defines the bipartite graph on edges "crossing the cut" i.e. edges with one endpoint in $\{v_1, \dots, v_i\}$ and the other endpoint in $\{v_{i+1}, \dots, v_n\}$. The maximum induced matching of G under layout σ is denoted $\text{mim}(\sigma, G)$, and is defined as the maximum, over all cuts of σ , of the value attained by the MIM of the cut, i.e. of the bipartite graph defined by the cut. The linear induced matching width – LMIM-width – of G is denoted $\text{lmw}(G)$, and is the minimum value of $\text{mim}(\sigma, G)$ over all possible linear orderings σ of the vertices of G .

We start by showing that if we have a path P in a tree T then the LMIM-width of T is no larger than the largest LMIM-width of any component of $T \setminus N[P]$, plus 1. To define these components the following notion is useful.

Definition 1 (Dangling tree). *Let T be a tree containing the adjacent nodes v and u . The dangling tree from v in u , $T\langle v, u \rangle$, is the component of $T \setminus (u, v)$ containing u .*

Given a node $x \in T$ with neighbours $\{v_1, \dots, v_d\}$, the forest obtained by removing $N[x]$ from T is a collection of dangling trees $\{T\langle v_i, u_{i,j} \rangle\}$, where $u_{i,j} \neq x$ is some neighbour of v_i . We can generalise this to a path $P = (x_1, \dots, x_p)$ in place of x , such that $T \setminus N[P] = \{T\langle v_{i,j}, u_{i,j,m} \rangle\}$, where $v_{i,j} \in N(P)$ is a neighbour of x_i and $u_{i,j,m} \notin N[P]$. See top part of Figure 1. This naming convention will be used in the following.

Lemma 1 (Path Layout Lemma). *Let T be a tree. If there exists a path $P = (x_1, \dots, x_p)$ in T such that every connected component of $T \setminus N[P]$ has*

LMIM-width $\leq k$ then $\text{lmw}(T) \leq k + 1$. Moreover, given the layouts for the components we can in linear time compute the layout for T .

Proof. Using the optimal linear orderings of the connected components of $T \setminus N[P]$, we give the below algorithm LINORD constructing a linear order σ_T on the nodes of T showing that $\text{lmw}(T)$ is $\leq k + 1$. The ordering σ_T starts out empty and the algorithm has an outer loop going through vertices in the path $P = (x_1, \dots, x_p)$. When arriving at x_i it uses the concatenation operator \oplus to add the path node x_i before looping over all neighbors $v_{i,j}$ of x_i adding the linear orders of each dangling tree from $v_{i,j}$ and then $v_{i,j}$ itself. See Figure 1 for an illustration.

```

function LINORD( $T$ : tree,  $P = (x_1, \dots, x_p)$ : path,  $\{\sigma_{T\langle v_{i,j}, u_{i,j,m} \rangle}$ : lin-ords)
   $\sigma_T \leftarrow \emptyset$  ▷ The list starts out empty
  for  $i \leftarrow 1, p$  do ▷ For all nodes on path  $(x_1, \dots, x_p)$ 
     $\sigma_T \leftarrow \sigma_T \oplus x_i$  ▷ Append path node
    for  $j \leftarrow 1, |N(x_i) \setminus P|$  do ▷ For all nbs of  $x_i$  not on path:  $v_{i,j}$ 
      for  $m \leftarrow 1, |N(v_{i,j}) \setminus x_i|$  do ▷ For all dangling trees from  $v_{i,j}$ 
         $\sigma_T \leftarrow \sigma_T \oplus \sigma_{T\langle v_{i,j}, u_{i,j,m} \rangle}$  ▷ Append given order of  $T\langle v_{i,j}, u_{i,j,m} \rangle$ 
       $\sigma_T \leftarrow \sigma_T \oplus v_{i,j}$  ▷ Append  $v_{i,j}$ 

```




Fig. 1. A tree with a path $P = (x_1, x_2, x_3, x_4)$, with nodes in $N[N[P]]$ and dangling trees featured, and below it the order given by the Path Layout Lemma

Firstly, from the algorithm it should be clear that each node of T is added exactly once to σ_T , that it runs in linear time, and that there is no cut containing two crossing edges from two separate dangling trees. Now we must show that σ_T does not contain cuts with MIM larger than $k + 1$. By assumption the layout of each dangling tree has no cut with MIM larger than k , and since these layouts can be found as subsequences of σ_T it follows that then also σ_T has no cut with more than k edges from a single dangling tree $T\langle v_{i,j}, u_{i,j,m} \rangle$. Also, we know that

edges from two separate dangling trees cannot both cross the same cut. The only edges of T left to account for, i.e. not belonging to one of the dangling trees, are those with both endpoints in $N[N[P]]$, the nodes at distance at most 2 from a node in P . For every cut of σ_T that contains more than a single crossing edge (x_i, x_{i+1}) there is a unique $x_i \in P$ and a unique $v_{i,j} \in N(x_i)$ such that every edge with both endpoints in $N[N[P]]$ that crosses the cut is incident on either x_i or $v_{i,j}$, and since the edge connecting x_i and $v_{i,j}$ also crosses the cut at most one of these edges can be taken into an induced matching. With these observations in mind, it is clear that $lmw(T) \leq mim(\sigma_T, T) \leq k + 1$.

Definition 2 (k -neighbour and k -component index). *Let x be a node in the tree T and v a neighbour of x . If v has a neighbour $u \neq x$ such that $lmw(T \setminus \langle v, u \rangle) \geq k$, then we call v a k -neighbour of x . The k -component index of x is equal to the number of k -neighbours of x and is denoted $D_T(x, k)$, or shortened to $D(x, k)$.*

Theorem 1 (Classification of LMIM-width of Trees). *For a tree T and $k \geq 1$ we have $lmw(T) \geq k + 1$ if and only if $D(x, k) \geq 3$ for some node x .*

Proof. We first prove the backward direction by contradiction. Thus we assume $D(x, k) \geq 3$ for a node x and there is a linear order σ such that $mim(\sigma, T) \leq k$.

Let v_1, v_2, v_3 be the three k -neighbors of x and T_1, T_2, T_3 the three trees of $T \setminus N[x]$ each of LMIM-width k , with v_i connected to a node of T_i for $i = 1, 2, 3$, that we know must exist by the definition of $D(x, k)$. We know that for each $i = 1, 2, 3$ we have a cut C_i in σ with $MIM=k$ and all k edges of this induced matching coming from the tree T_i . Wlog we assume these three cuts come in the order C_1, C_2, C_3 , i.e. with the cut having an induced matching of k edges of T_2 in the middle. Note that in σ all nodes of T_1 must appear before C_2 and all nodes of T_3 after C_2 , as otherwise, since T is connected and the distance between T_2 and the two trees T_1 and T_3 is at least two, there would be an extra edge crossing C_2 that would increase MIM of this cut to $k + 1$. It is also clear that v_1 has to be placed before C_2 and v_3 has to be placed after C_2 , for the same reason, e.g. the edge between v_1 and a node of T_1 cannot cross C_2 without increasing MIM. But then we are left with the vertex x that cannot be placed neither before C_2 nor after C_2 without increasing MIM of this cut by adding at least one of (v_1, x) or (v_3, x) to the induced matching. We conclude that $D(x, k) \geq 3$ for a node x implies LMIM-width at least $k + 1$.

To prove the forward direction we first show the following partial claim: if $lmw(T) \geq k + 1$ then there exists a node $x \in T$ such that $D(x, k) \geq 3$; or there exists a strict subtree S of T with $lmw(S) \geq k + 1$. We will prove the contrapositive statement, so let us assume that every node in T has $D(x, k) < 3$ and no strict subtree of T has LMIM-width $\geq k + 1$ and show that then $lmw(T) \leq k$. For every node $x \in T$, it must then be true that $D(x, k) \leq 2$ and that $D(x, k + 1) = 0$. The strategy of this proof is to show that there is always a path P in T such that all the connected components in $T \setminus N[P]$ have LMIM-width $\leq k - 1$. When we have shown this, we proceed to use the Path

Layout Lemma, to get that $lmw(T) \leq k$. To prove this, we define the following two sets of vertices:

$$X = \{x | x \in V(T) \text{ and } D(x, k) = 2\}, Y = \{y | y \in V(T) \text{ and } D(y, k) = 1\}$$

Case 1: $X \neq \emptyset$

If x_i and x_j are in X , then every vertex on the path $P(x_i, \dots, x_j)$ connecting x_i and x_j must be elements of X , as every node on this path clearly has a dangling tree with LMIM-width k in the direction of x_i and in the direction of x_j . The fact that every pair of vertices in X are connected by a path in X means that X must be a connected subtree of T . Furthermore, this subtree must be a path, otherwise there are three disjoint dangling trees $T\langle v_1, u_1 \rangle, T\langle v_2, u_2 \rangle, T\langle v_3, u_3 \rangle$, each with LMIM-width k , and each hanging from a separate node. But then there is some vertex w such that $T\langle v_1, u_1 \rangle, T\langle v_2, u_2 \rangle$ and $T\langle v_3, u_3 \rangle$ are subtrees of dangling trees from different neighbours of w . But this implies that $D(w, k) \geq 3$, which we assumed were not the case, so this leads to a contradiction. We therefore conclude that all nodes in X must lie on some path $P = (x_1, \dots, x_p)$. The final part of the argument lies in showing that we can apply the Path Layout Lemma. For some $x_i \in P, i \in \{2, \dots, p-1\}$, its k -neighbours are x_{i-1} and x_{i+1} . For x_1 , these neighbours are x_2 and some $x_0 \notin X$. For x_p , these neighbours are x_{p-1} and some $x_{p+1} \notin X$. x_0 and x_{p+1} may only have one k -neighbour – x_1 and x_p respectively – or else they would be in X . If we make $P' = (x_0, \dots, x_{p+1})$, we then see that every connected component in $T \setminus N[P']$ must have LMIM-width $\leq k-1$. By the Path Layout Lemma, $lmw(T) \leq k$.

Case 2: $X = \emptyset, Y \neq \emptyset$

We construct the path P in a simple greedy manner as follows. We start with $P = (y_1, y_2)$, where y_1 is some arbitrary node in Y , and y_2 its only k -neighbour. Then, if the highest-numbered node in P , call it y_q , has a k -neighbour $y' \notin P$, then we assign y_{q+1} to y' , and repeat this process exhaustively. Since we look at finite graphs, we will eventually reach some node y_p such that either $y_p \notin Y$ or y_p 's k -neighbour is y_{p-1} . We are then done and have $P = (y_1, \dots, y_p)$, which must be a path in T , since every node $y_{i+1} \in P$ is a neighbour of y_i and for y_i we only assign maximally one such y_{i+1} . Also, every connected component of $T \setminus N[P]$ must have LMIM-width $\leq k-1$. If not, some node $y_i \in P$ would have a k -neighbour $y' \notin P$, but by the assumption $X = \emptyset$ this is impossible, since then either $i < p$ and y_i has two k -neighbours y' and y_{i+1} , or else $i = p$ and $y_p \in Y$ and y_i has the two k -neighbors y' and y_{i-1} (in case $i = p$ and $y_p \notin Y$ then by definition of Y the node y_i could not have a k -neighbor y'). By the Path Layout Lemma, $lmw(T) \leq k$.

Case 3: $X = \emptyset, Y = \emptyset$

If you make $P = (x)$ for some arbitrary $x \in T$, it is obvious that every connected component of $T \setminus N[P]$ has LMIM-width $\leq k-1$. By the Path Layout Lemma, $lmw(T) \leq k$.

We have proven the partial claim that if $lmw(T) \geq k + 1$ then there exists a node $x \in T$ such that $D(x, k) \geq 3$; or there exists a strict subtree S of T with $lmw(S) \geq k + 1$. To finish the backward direction of the theorem we need to show that if $lmw(T) \geq k + 1$ then there exists a node $x \in T$ with $D(x, k) \geq 3$. Assume for a contradiction that there is no node with k -component index at least 3 in T . By the partial claim, there must then exist a strict subtree S with $lmw(S) \geq k + 1$. But since we look at finite trees, we know that there in S must exist a minimal subtree S_0 , $lmw(S_0) = k + 1$ with no strict subtree with LMIM-width $> k$. By the partial claim, S_0 must contain a node x_0 with $D_{S_0}(x_0, k) \geq 3$. But every dangling tree $S_0(v, u)$ is a subtree of $T(v, u)$, and so if $D_{S_0}(x_0, k) \geq 3$, then $D_T(x_0, k) \geq 3$ contradicting our assumption.




Fig. 2. The smallest tree with LMIM-width 2, having a node v with three 1-neighbors u_1, u_2, u_3 having dangling trees S_1, S_2, S_3 , respectively, so that $D(v, 1) = 3$

By Theorem 1, every tree with LMIM-width $k \geq 2$ must be at least 3 times bigger than the smallest tree with LMIM-width $k - 1$, which implies the following.

Remark 1. The LMIM-width of an n -node tree is $\mathcal{O}(\log n)$.

3 Rooted trees, k -critical nodes and labels

Our algorithm computing LMIM-width will work on a rooted tree, processing it bottom-up. We will choose an arbitrary node r of the tree T and denote by T_r the tree rooted in r . For any node x we denote by $T_r[x]$ the standard *complete subtree* of T_r rooted in x . During the bottom-up processing of T_r we will compute a label for various subtrees. The notion of a k -critical node is crucial for the definition of labels.

Definition 3 (k -critical node). Let T_r be a rooted tree with $lmw(T_r) = k$. We call a node x in T_r *k -critical* if it has exactly two children v_1 and v_2 that each has at least one child, u_1 and u_2 respectively, such that $lmw(T_r[u_1]) = lmw(T_r[u_2]) = k$. Thus x is k -critical if and only if $lmw(T) = k$ and $D_{T_r[x]}(x, k) = 2$.

Remark 2. If T_r has LMIM-width k it has at most one k -critical node.

Proof. For a contradiction, let x and x' be two k -critical nodes in T_r . There are then four nodes, v_l, v_r, v'_l, v'_r , the two k -neighbours of x and x' respectively, such that there exist dangling trees $T\langle v_l, u_l \rangle, T\langle v_r, u_r \rangle, T\langle v'_l, u'_l \rangle, T\langle v'_r, u'_r \rangle$ that all have LMIM-width k . If x and x' have a descendant/ancestor relationship in T_r , then assume wlog that x' is a descendant of v_l , and note that $T\langle v_r, u_r \rangle, T\langle v'_l, u'_l \rangle$ and $T\langle v'_r, u'_r \rangle$ are disjoint trees in different neighbours of x' , thus $D_{T_r}(x', k) = 3$ and by Theorem 1 T_r should have LMIM-width $k + 1$. Otherwise, all the dangling trees are disjoint, thus $D_T(x, k) = D_T(x', k) = 3$ and we arrive at the same conclusion.

Definition 4 (label). Let rooted tree T_r have $lmw(T_r) = k$. Then $\mathbf{label}(T_r)$ consists of a list of decreasing numbers, (a_1, \dots, a_p) , where $a_1 = k$, appended with a string called *last_type*, which tells us where in the tree an a_p -critical node lies, if it exists at all. If $p = 1$ then the label is simple, otherwise it is complex. The $\mathbf{label}(T_r)$ is defined recursively, with type 0 being a base case for singletons and for stars, and with type 4 being the only one defining a complex label.

- Type 0: r is a leaf, i.e. T_r is a singleton, then $\mathbf{label}(T_r) = (0, t.0)$;
or all children of r are leaves, then $\mathbf{label}(T_r) = (1, t.0)$
- Type 1: No k -critical node in T_r , then $\mathbf{label}(T_r) = (k, t.1)$
- Type 2: r is the k -critical node in T_r , then $\mathbf{label}(T_r) = (k, t.2)$
- Type 3: A child of r is k -critical in T_r , then $\mathbf{label}(T_r) = (k, t.3)$
- Type 4: There is a k -critical node u_k in T_r that is neither r nor a child of r .
Let w be the parent of u_k . Then $\mathbf{label}(T_r) = k \oplus \mathbf{label}(T_r \setminus T_r[w])$

In type 4 we note that $lmw(T_r \setminus T_r[w]) < k$ since otherwise u_k would have three k -neighbors (two children in the tree and also its parent) and by Theorem 1 we would then have $lmw(T_r) = k + 1$. Therefore, all numbers in $\mathbf{label}(T_r \setminus T_r[w])$ are smaller than k and a complex label is a list of decreasing numbers followed by *last_type* $\in \{t.0, t.1, t.2, t.3\}$. We now give a Proposition that for any node x in T_r will be used to compute $\mathbf{label}(T_r[x])$ based on the labels of the subtrees rooted at the children and grand-children of x . The subroutine underlying this Proposition, see the decision tree in Figure 3, will be used when reaching node x in the bottom-up processing of T_r .

Proposition 1. Let x be a node of T_r with children $Child(x)$, and given $\mathbf{label}(T_r[v])$ for all $v \in Child(x)$. We define (and compute) $k = \max_{v \in Child(x)} \{lmw(T_r[v])\}$ and $N_k = \{v \in Child(x) \mid lmw(T_r[v]) = k\}$ and denote by $N_k = \{v_1, \dots, v_q\}$ and by $l_i = \mathbf{label}(T_r[v_i])$. Define (compute) $t_k = D_{T_r[x]}(x, k)$ by noting that $t_k = |\{v_i \in N_k \mid v_i \text{ has child } u_j \text{ with } lmw(T_r[u_j]) = k\}|$. Given this information, we can find $\mathbf{label}(T_r[x])$ as follows:

- **Case 0:** if $|Child(x)| = 0$ then $\mathbf{label}(T_r[x]) = (0, t.0)$;
else if $k = 0$ then $\mathbf{label}(T_r[x]) = (1, t.0)$
- **Case 1:** Every label in N_k is simple and has *last_type* equal to $t.1$ or $t.0$, and $t_k \leq 1$. Then, $\mathbf{label}(T_r[x]) = (k, t.1)$

- **Case 2:** Every label in N_k is simple and has *last_type* equal to *t.1* or *t.0*, but $t_k = 2$. Then, $\text{label}(T_r[x]) = (k, t.2)$
- **Case 3:** Every label in N_k is simple and has *last_type* equal to *t.1* or *t.0*, but $t_k \geq 3$. Then, $\text{label}(T_r[x]) = (k + 1, t.1)$
- **Case 4:** $|N_k| \geq 2$ and for some $v_i \in N_k$, either l_i is a complex label, or l_i has *last_type* equal to either *t.2* or *t.3*. Then, $\text{label}(T_r[x]) = (k + 1, t.1)$
- **Case 5:** $|N_k| = 1$, l_1 is a simple label and l_1 has *last_type* equal to *t.2*. Then, $\text{label}(T_r[x]) = (k, t.3)$
- **Case 6:** $|N_k| = 1$, l_1 is either complex or has *last_type* equal to *t.3*, and $k \notin \text{label}(T_r[x] \setminus T_r[w])$, where w is the parent of the k -critical node in $T_r[v_1]$. Then, $\text{label}(T_r[x]) = k \oplus \text{label}(T_r[x] \setminus T_r[w])$
- **Case 7:** $|N_k| = 1$, l_1 is either complex or has *last_type* equal to *t.3*, and $k \in \text{label}(T_r[x] \setminus T_r[w])$, where w is the parent of the k -critical node in $T_r[v_1]$. Then, $\text{label}(T_r[x]) = (k + 1, t.1)$

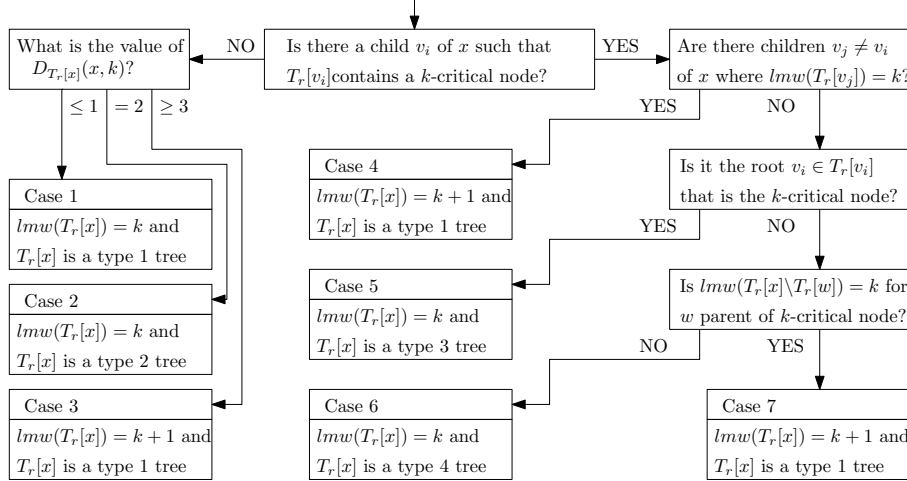


Fig. 3. A decision tree corresponding to the case analysis of Proposition 1

Proof. We show that exactly one case applies to every rooted tree and in each case we assign the label according to Definition 4. First the base case: either x is a leaf or all its children are leaves and we are in Case 0 and the label is assigned according to Def. 4. Otherwise, observe the decision tree in Figure 3. It follows from Def. 4, k , N_k and t_k that cases 1 up to 7 of Prop. 1 corresponds to cases 1 up to 7 in the decision tree - we mention this correspondence in the below - and this proves that exactly one case applies to every rooted tree. The following facts simplify the case analysis: $\text{lmw}(T_r[x])$ is equal to either k or $k + 1$, and since no subtree rooted in a child of x has LMIM-width $k + 1$ there cannot be any $(k + 1)$ -critical node in $T_r[x]$, therefore if $\text{lmw}(T_r[x]) = k + 1$, $T_r[x]$ is always a type 1

tree and by Theorem 1 it must contain a node v such that $D_{T_r[x]}(v, k) \geq 3$. This node must either be a k -critical node in a rooted subtree of $T_r[x]$, or x itself. We go through the cases 1 to 7 in order.

Note that in Cases 1, 2, and 3 the condition 'Every label in N_k is simple and has *last_type* equal to $t.1$ or $t.0$ ' means there are no k -critical nodes in any subtree of $T_r[x]$, because every $T_r[v]$ for $v \in \text{Child}(x)$ is either of type 1 or has LMIM-width $< k$:

Case 1: By definition of t_k , $D_{T_r[x]}(x, k) \leq 1$. Therefore, $lmw(T_r[x]) = k$, and $T_r[x]$ is a type 1 tree.

Case 2: By definition of t_k , $D_{T_r[x]}(x, k) = 2$, and no other nodes are k -critical, therefore $lmw(T_r[x]) = k$. But now x is k -critical in $T_r[x]$ so $T_r[x]$ is a type 2 tree.

Case 3: By definition of t_k , $D_{T_r[x]}(x, k) = 3$ and $lmw(T_r[x]) = k + 1$.

For the remaining Cases 4, 5, 6 and 7, some $T_r[v]$ for $v \in \text{Child}(x)$ has LMIM-width k and is of type 2, 3 or 4, so at least one k -critical node exists in some subtree of $T_r[x]$:

Case 4: There is a k -critical node u_k in some $T_r[v_i]$ (not of type 1), and some other v_j has $lmw(T_r[v_j]) = k$ (because $|N_k| \geq 2$). Now observe w the parent of u_k . The dangling tree $T_r[x] \setminus T_r[w]$ is a supertree of $T_r[v_j]$ and thus has LMIM-width $\geq k$. Therefore w is a k -neighbour of u_k and by Theorem 1 $lmw(T_r[x]) = k + 1$.

Case 5: x has only one child v with $lmw(T_r[v]) = k$, and v is itself k -critical ($T_r[v]$ is type 2). x cannot be a k -neighbour of v in the unrooted $T_r[x]$, because every dangling tree from x is some $T_r[v_i]$, $v_i \neq v$ of x , which we know has LMIM-width $< k$. Since no other node in T is k -critical, $lmw(T_r[x]) = k$, and since v , a child of x , is k -critical in $T_r[x]$, $T_r[x]$ is a type 3 tree.

Case 6: x has only one child v with $lmw(T_r[v]) = k$, and there is a k -critical node u_k with parent w – neither of which are equal to x – in $T_r[v]$ ($T_r[v]$ is a type 3 or type 4 tree). Moreover, no tree rooted in another child of w , apart from u_k , can have LMIM-width $\geq k$, since this would imply $D_{T_r[v]}(u_k, k) = 3$ and thus $lmw(T_r[v]) > k$; nor can $T_r[x] \setminus T_r[w]$ have LMIM-width $= k$, since then we would have k in $label(T_r[x] \setminus T_r[w])$ disagreeing with the condition of Case 6. Therefore $D_{T_r[x]}(u, k) = 2$, and $lmw(T_r[x]) = k$. $T_r[x]$ is thus a type 4 tree and the label is assigned according to the definition.

Case 7: $T_r[v]$, u_k and w are as described in Case 6. But here, $lmw(T_r[x] \setminus T_r[w]) = k$ (since the condition says that k is in its label), and thus w is a k -neighbour of its child u_k and by Theorem 1 $lmw(T_r[x]) = k + 1$.

We conclude that $label(T_r[x])$ has been assigned the correct value in all possible cases.

4 Computing LMIM-width of Trees and Finding a Layout

The subroutine underlying Prop. 1 will be used in a bottom-up algorithm that starts out at the leaves and works its way up to the root, computing labels

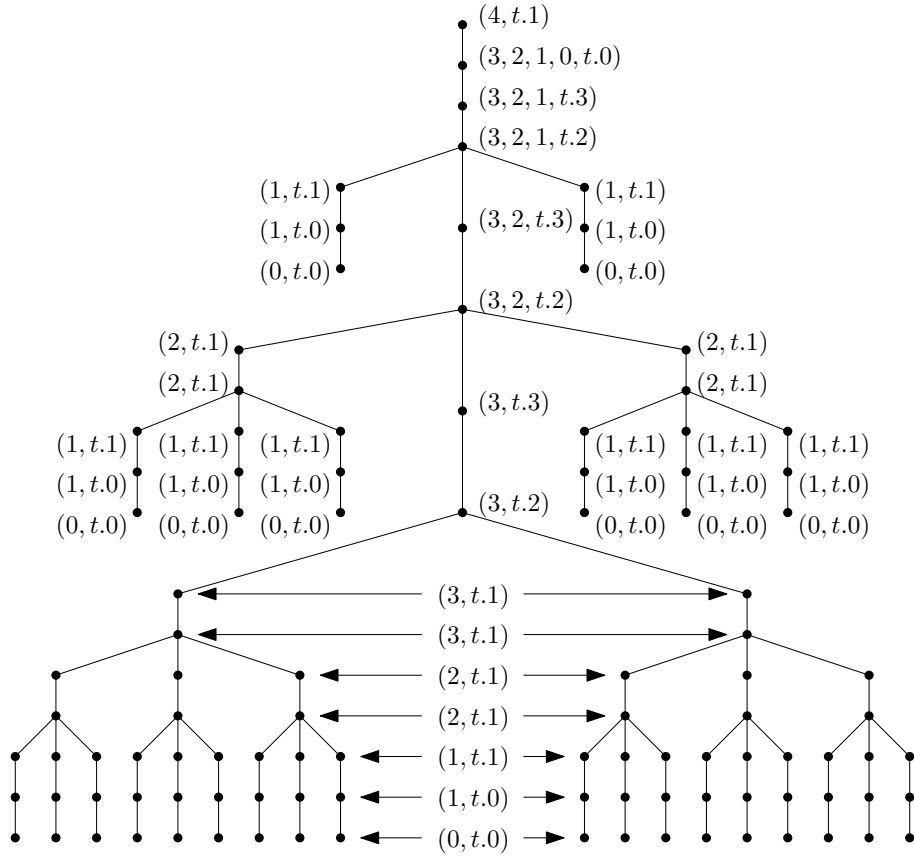


Fig. 4. A rooted tree of LMIM-width 4 with labels of subtrees. We explain the labels $(3, t.2)$, $(3, t.3)$, $(3, 2, t.2)$ assigned to subtrees rooted at the nodes we call a, b, c , with $\text{parent}(a) = b$ and $\text{parent}(b) = c$. The sub-tree rooted at a , with label $(3, t.2)$ has precisely two children that have a child-tree each of LMIM-width 3, hence a is 3-critical and it is a type 2 tree (Case 2 of Prop. 1). The sub-tree rooted at b , labelled $(3, t.3)$, is thus the parent of a 3-critical node, and so it is of type 3 (Case 5 of Prop. 1). The sub-tree rooted at c with label $(3, 2, t.2)$ has maximum LMIM-width of a child-tree being 3, and it has a 3-critical node a which is neither c nor a child of c , so it is of type 4 (Case 6 of Prop. 1); and moreover the subtree $T_r[c] \setminus T_r[a]$ has LMIM-width 2 with node c as 2-critical so it is of type 2 (Case 2 of Prop. 1), and the label of $T_r[c]$ becomes $3 \oplus (2, t.2)$.

of subtrees $T_r[x]$. However, in two cases (Case 6 and 7) we need the label of $T_r[x] \setminus T_r[w]$, which is not a complete subtree rooted in any node of T_r . Note that the label of $T_r[x] \setminus T_r[w]$ is again given by a (recursive) call to Prop. 1 and is then stored as a suffix of the complex label of $T_r[x]$. We will compute these labels by iteratively calling Prop. 1 (substituting the recursion by iteration). We first need to carefully define the subtrees involved when dealing with complex labels.

From the definition of labels it is clear that only type 4 trees lead to a complex label. In that case we have a tree $T_r[x]$ of LMIM-width k and a k -critical node u_k that is neither x nor a child of x , and the recursive definition gives $label(T_r[x]) = k \oplus label(T_r[x] \setminus T_r[w])$ for w the parent of u_k . Unravelling this recursive definition, this means that if $label(T_r[x]) = (a_1, \dots, a_p, last_type)$, we can define a list of nodes (w_1, \dots, w_{p-1}) where w_i is the parent of an a_i -critical node in $T_r[x] \setminus (T_r[w_1] \cup \dots \cup T_r[w_{i-1}])$. We expand this list with $w_p = x$, such that there is one node in $T_r[x]$ corresponding to each number in $label(T_r[x])$, and $T_r[x] \setminus (T_r[w_1] \cup \dots \cup T_r[w_p]) = \emptyset$.

Now, in the first level of a recursive call to Prop. 1 the role of $T_r[x]$ is taken by $T_r[x] \setminus T_r[w_1]$, and in the next level it is taken by $(T_r[x] \setminus T_r[w_1]) \setminus T_r[w_2]$ etc. The following definition gives a shorthand for denoting these trees.

Definition 5. Let x be a node in T_r , $label(T_r[x]) = (a_1, a_2, \dots, a_p, last_type)$ and the corresponding list of vertices (w_1, \dots, w_p) is as we describe in the above text. For any non-negative integer s , the tree $\mathbf{T}_r[\mathbf{x}, \mathbf{s}]$ is the subtree of $T_r[x]$ obtained by removing all trees $T_r[w_i]$ from $T_r[x]$, where $a_i \geq s$. In other words, if q is such that $a_q \geq s > a_{q+1}$, then $T_r[x, s] = T_r[x] \setminus (T_r[w_1] \cup T_r[w_2] \cup \dots \cup T_r[w_q])$

Remark 3. Some important properties of $T_r[x, s]$ are the following. Let $T_r[x, s]$, $label(T_r[x, s])$, (w_1, \dots, w_p) and q as in the definition. Then

1. if $s > a_1$, then $T_r[x, s] = T_r[x]$
2. $label(T_r[x, s]) = (a_{q+1}, \dots, a_p, last_type)$
3. $lmw(T_r[x, s]) = a_{q+1} < s$
4. $lmw(T_r[x, s+1]) = s$ if and only if $s \in label(T_r[x])$
5. $T_r[x, s+1] \neq T_r[x, s]$ if and only if $s \in label(T_r[x])$

Proof. These follow from the definitions, maybe the last one requires a proof:

Backward direction: Let $s = a_q$ for some $1 \leq q \leq p$. Then $T_r[x, s+1] = T_r[x] \setminus (T_r[w_1] \cup \dots \cup T_r[w_{q-1}])$ and $T_r[x, s] = T_r[x] \setminus (T_r[w_1] \cup \dots \cup T_r[w_q])$. These two trees are clearly different.

Forward direction: Let $T_r[x, s] = T_r[x] \setminus (T_r[w_1] \cup \dots \cup T_r[w_q])$ and $T_r[x, s+1] = T_r[x] \setminus (T_r[w_1] \cup \dots \cup T_r[w_{q'}])$ with $q' < q$ and $a_{q'} > a_q$ (because numbers in a label are strictly descending). $a_q < s+1$ and $a_q \geq s$, ergo $a_q = s$.

Note that for any s the tree $T_r[x, s]$ is defined only after we know $label(T_r[x])$. In the algorithm, we compute $label(T_r[x])$ by iterating over increasing values of s (until $s > lmw(T_r[x])$) since by Remark 3.1 we then have $T_r[x, s] = T_r[x]$ and we could hope for a loop invariant saying that we have correctly computed $label(T_r[x, s])$. However, $T_r[x, s]$ is only known once we are done. Instead, each iteration of the loop will correctly compute the label of the following subtree called $T_{union}[x, s]$, which is not always equal to $T_r[x]$, but importantly for $s > lmw(T_r[x])$, we will have $T_{union}[x, s] = T_r[x, s] = T_r[x]$.

Definition 6. Let x be a node in T_r with children v_1, \dots, v_d . $T_{union}[x, s]$ is then equal to the tree induced by x and the union of all $T_r[v_i, s]$ for $1 \leq i \leq d$. More technically, $T_{union}[x, s] = T_r[V']$ where $V' = x \cup V(T_r[v_1, s]) \cup \dots \cup V(T_r[v_d, s])$.

Given a tree T , we find its LMIM-width by rooting it in an arbitrary node r , and computing labels by processing T_r bottom-up. The answer is given by the first element of $label(T_r[r])$, which by definition is equal to $lmw(T)$. At a leaf x of T_r we initialize by $label(T_r[x]) \leftarrow (0, t.0)$, and at a node x for which all children are leaves we initialize by $label(T_r[x]) \leftarrow (1, t.0)$, according to Definition 4. When reaching a higher node x we compute label of $T_r[x]$ by calling function $MAKELABEL(T_r, x)$.

```

function MAKELABEL( $T_r, x$ )                                 $\triangleright$  finds  $cur\_label = label(T_r[x])$ 
     $cur\_label \leftarrow (0, t.0)$                                  $\triangleright$  This is  $label(T_{union}[x, 0])$ 
     $\{v_1, \dots, v_d\} = \text{children of } x$ 
    if  $0 \in label(T_r[v_i])$  for some  $i$  then
         $cur\_label \leftarrow (1, t.0)$                                  $\triangleright$  This is then  $label(T_{union}[x, 1])$ 
    for  $s \leftarrow 1, \max_{i=1}^d \{\text{first element of } label(T_r[v_i])\}$  do
         $\{l'_1, \dots, l'_d\} = \{label(T_r[v_i, s+1]) \mid 1 \leq i \leq d\}$ 
         $N_s = \{v_i \mid 1 \leq i \leq d, s \in l'_i\}$ 
         $t_s = |\{v_i \mid v_i \in N_s, v_i \text{ has child } u_j \text{ s.t. } s \in label(T_r[u_j, s+1])\}|$ 
        if  $|N_s| > 0$  then
             $case \leftarrow$  the case from Prop. 1 applying to  $s, \{l'_1, \dots, l'_d\}, N_s$  and  $t_s$ 
             $cur\_label \leftarrow$  as given by  $case$  in Prop. 1 ( $s \oplus cur\_label$  if Case 6)
    
```

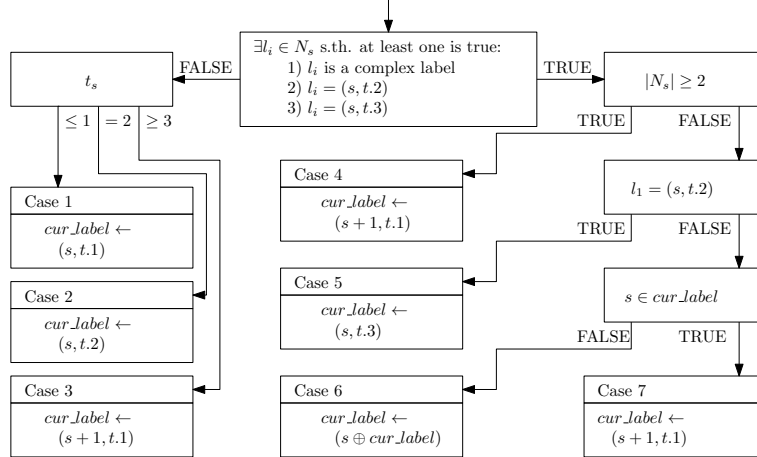


Fig. 5. The same decision tree as shown in Prop. 1, but adapted to MAKELABEL

Lemma 2. Given labels at descendants of node x in T_r , $MAKELABEL(T_r, x)$ computes $label(T_r[x])$ as the value of cur_label .

Proof. Assume that x has the children v_1, \dots, v_d , and denote their set of labels as $L = \{l_1, \dots, l_d\}$. MAKELABEL keeps a variable cur_label that is updated

maximally k times in a for loop, where k is the biggest number in any label of children of x . The following claim will suffice to prove the lemma, since for $s > \text{lmw}(T_r[x])$, we have $T_{\text{union}}[x, s] = T_r[x]$.

Claim: At the end of the s 'th iteration of the for loop the value of cur_label is equal to $\text{label}(T_{\text{union}}[x, s+1])$.

Base case: We have to show that before the first iteration of the loop we have $\text{cur_label} = \text{label}(T_{\text{union}}[x, 1])$. If some label $l_i \in L$ has 0 as an element then $T_{\text{union}}[x, 1]$ is isomorphic to a star with x as the center and v_i as a leaf. By Prop. 1, in this case $\text{label}(T_{\text{union}}[x, 1]) = (1, t.0)$ and this is what cur_label is initialized to. If no $l_i \in L$ has 0 as an element, then by Remark 3.5 $T_{\text{union}}[x, 1] = T_{\text{union}}[x, 0]$ which by definition is the singleton node x and by Prop. 1 the label of this tree is $(0, t.0)$ and this is what cur_label is initialized to.

Induction step: We assume $\text{cur_label} = \text{label}(T_{\text{union}}[x, s])$ at the start of the s 'th iteration of the for loop and show that at the end of the iteration, $\text{cur_label} = \text{label}(T_{\text{union}}[x, s+1])$.

The first thing done in the for loop is the computation of $\{l'_i \mid 1 \leq i \leq d, l'_i = \text{label}(T_r[v_i, s+1])\}$. By Remark 3.2, $\text{label}(T_r[v_i, s+1]) \subseteq \text{label}(T_r[v_i])$ for all i , therefore l'_1, \dots, l'_d are trivial to compute. The second thing done is to set N_s as the set of all children of x whose labels contain s , and t_s as the number of nodes in N_s that themselves have children whose labels contain s . Let us first look at what happens when $|N_s| = 0$:

By Remark 3.5, for every child v_i of x , $T_r[v_i, s+1] = T_r[v_i, s]$ if $s \notin \text{label}(T_r[v_i])$. Therefore, if $|N_s| = 0$, then $T_{\text{union}}[x, s+1] = T_{\text{union}}[x, s]$, and from the induction assumption, $\text{label}(T_{\text{union}}[x, s+1]) = \text{cur_label}$, and indeed when $|N_s| = 0$ then iteration s of the loop does not alter cur_label .

Otherwise, we have $|N_s| > 0$ and make a call to the subroutine given by Prop. 1, see the decision tree in Figure 5, to compute $\text{label}(T_{\text{union}}[x, s+1])$ and argue first that the variables used in that call correspond to the variables used in Prop. 1 to compute $\text{label}(T_r[x])$. The correspondence is given in Table 4. Most of these are just observations: $T_{\text{union}}[x, s+1]$ corresponds to $T_r[x]$

Proposition 1	for loop iteration s	Explanation
$T_r[x], k$	$T_{\text{union}}[x, s+1], s$	Tree needing label, max lmw of children
$T_r[v_1], \dots, T_r[v_d]$	$T_r[v_i, s], \dots, T_r[v_d, s]$	Subtrees of children
$l_1, \dots, l_d, N_k, t_k$	$l'_1, \dots, l'_d, N_s, t_s$	Child labels, those with max, root comp. index
$\text{label}(T_r[x] \setminus T_r[w])$	cur_label	This is also $\text{label}(T_{\text{union}}[x, s+1] \setminus T_r[w, s+1])$

in Prop. 1, and $T_r[v_1, s+1], \dots, T_r[v_d, s+1]$ corresponds to $T_r[v_1], \dots, T_r[v_d]$. $\{l'_i \mid 1 \leq i \leq d, l'_i = \text{label}(T_r[v_i, s+1])\}$ correspond to $\{\text{label}(T_r[v]) \mid v \in \text{Child}\}$ in Prop. 1. N_s is defined in the algorithm so that it corresponds to N_k in Prop. 1. Since $|N_s| > 0$, some v_i has s in its label l'_i . By Remark 3.3 and 3.4, we can infer that s is the maximum LMIM-width of all $T_r[v_i, s+1]$, therefore s corresponds

to k in Proposition 1.

It takes a bit more effort to show that t_s computed in iteration s of the for loop corresponds to $t_k = D_{T_r[x]}(x, k)$ in Prop. 1 – meaning we need to show that $t_s = D_{T_{union}[x, s+1]}(x, s)$. Consider v_i , a child of x . In accordance with MAKELABEL we say that v_i contributes to t_s if $v_i \in N_s$ and v_i has a child u_j with s in its label. We thus need to show that v_i contributes to t_s if and only if v_i is an s -neighbour of x in $T_{union}[x, s+1]$. Observe that by Remark 3.4, $lmw(T_r[v_i, s+1]) = lmw(T_r[u_j, s+1]) = s$ if and only if s is in the labels of both $T_r[v_i]$ and $T_r[u_j]$. If $s \notin \text{label}(T_r[u_j, s+1])$, then $lmw(T_r[u_j, s+1]) < s$, and if this is true for all children of v_i , then v_i is not an s -neighbour of x in $T_{union}[x, s+1]$. If $s \notin \text{label}(T_r[v_i, s+1])$, then $lmw(T_r[v_i, s+1]) < s$ and no subtree of $T_r[v_i, s+1]$ can have LMIM-width s . However, if $s \in \text{label}(T_r[u_j, s+1])$ and $s \in \text{label}(T_r[v_i, s+1])$ (this is when v_i contributes to t_s), then $T_r[v_i, s+1] \cap T_r[u_j]$ must be equal to $T_r[u_j, s+1]$ and $T_r[u_j, s+1] \subseteq T_{union}[x, s+1]$, and we conclude that v_i is an s -neighbour of x in $T_{union}[x, s+1]$ if and only if v_i contributes to t_s , so $t_s = D_{T_{union}[x, s+1]}(x, s)$.

Lastly, we show that if $T_{union}[x, s+1]$ is a Case 6 or Case 7 tree – that is, $|N_s| = 1$, and $T_r[v_1, s+1]$ is a type 3 or type 4 tree, with w being the parent of an s -critical node – then the algorithm has $\text{label}(T_{union}[x, s+1] \setminus T_r[w, s+1])$ available for computation, indeed that this is the value of cur_label . We know, by definition of label and Remark 3.5 that $T_r[v_i, s+1] \setminus T_r[v_i, s] = T_r[w, s+1]$. But since $|N_s| = 1$, for every $j \neq i$, $T_r[v_j, s+1] \setminus T_r[v_j, s] = \emptyset$. Therefore $T_{union}[x, s+1] \setminus T_{union}[x, s] = T_r[w, s+1]$ and $T_{union}[x, s+1] \setminus T_r[w, s+1] = T_{union}[x, s]$. But by the induction assumption, $\text{cur_label} = \text{label}(T_{union}[x, s])$. Thus cur_label corresponds to $\text{label}(T_r[x] \setminus T_r[w])$ in Prop. 1.

We have now argued for all the correspondences in Table 4. By that, we conclude from Prop. 1 and Definition ?? and the inductive assumption that $\text{cur_label} = \text{label}(T_{union}[x, s+1])$ at the end of the s 'th iteration of the for loop in MAKELABEL. It runs for k iterations, where k is equal to the biggest number in any label of the children of x , and cur_label is then equal to $\text{label}(T_{union}[x, k+1])$. Since $k \geq lmw(T_r[v_i])$ for all i , by definition $T_r[v_i, k+1] = T_r[v_i]$ for all i , and thus $T_{union}[x, k+1] = T_r[x]$. Therefore, when MAKELABEL finishes, $\text{cur_label} = \text{label}(T_r[x])$.

Theorem 2. *Given any tree T , $lmw(T)$ can be computed in $\mathcal{O}(n \log(n))$ -time.*

Proof. We find $lmw(T)$ by bottom-up processing of T_r and returning the first element of $\text{label}(T_r)$. After correctly initializing at leaves and nodes whose children are all leaves, we make a call to MAKELABEL for each of the remaining nodes. Correctness follows by Lemma 2 and induction on the structure of the rooted tree. For the timing we show that each call runs in $\mathcal{O}(\log n)$ time. For every integer s from 1 to m , the biggest number in any label of children of x , which is $\mathcal{O}(\log n)$ by Remark 1, the algorithm checks how many labels of children of x contain s (to compute N_s), and how many labels of grandchildren of x contain s (to compute t_s). The labels are sorted in descending order, therefore the whole loop goes only once through each of these labels, each of length

$O(\log n)$. Other than this, MAKELABEL only does a constant amount of work. Therefore, $\text{MAKELABEL}(T_r, x)$, if x has a children and b grandchildren, takes time proportional to $O(\log n)(a + b)$. As the sum of the number of children and grandchildren over all nodes of T_r is $O(n)$ we conclude that the total runtime to compute $lmw(T)$ is $\mathcal{O}(n \cdot \log n)$.

Theorem 3. *A layout of LMIM-width $lmw(T)$ of a tree T can be found in $\mathcal{O}(n \cdot \log n)$ -time.*

Proof. Given T we first run the algorithm computing $lmw(T)$ by finding labels of all nodes and various subtrees. Given T we first run the algorithm computing $lmw(T)$ finding the label of every full rooted subtree in T_r . We give a recursive layout-algorithm that uses these labels in tandem with LINORD presented in the Path Layout Lemma. We call it on a rooted tree where labels of all subtrees are known. For simplicity we call this rooted tree T_r even though in recursive calls this is not the original root r and tree T . The layout-algorithm goes as follows:

- 1) Let $lmw(T_r) = k$ and find a path P in T_r such that all trees in $T_r \setminus N[P]$ have LMIM-width $< k$. The path depends on the type of T_r as explained in detail below.
- 2) Call this layout-algorithm recursively on every rooted tree in $T_r \setminus N[P]$ to obtain linear layouts; to this end, we need the correct label for every node in these trees.
- 3) Call LINORD on T_r , P and the layouts provided in step 2.

Every tree in the forest $T \setminus N[P]$ is equal to a dangling tree $T\langle v, u \rangle$, where v is a neighbour of some $x \in P$.

We observe that if $lmw(T) = k$, then by definition $lmw(T\langle v, u \rangle) = k$ if and only if v is a k -neighbour of x . It follows that every tree in $T \setminus N[P]$ has LMIM-width at most $k - 1$ if and only if no node in P has a k -neighbour that is not in P . We use this fact to show that for every type of tree we can find a satisfying path in the following way:

Type 0 trees: Choose $P = (r)$. Since $T \setminus N[r] = \emptyset$ in these trees, this must be a satisfying path.

Type 1 trees: These trees contain no k -critical nodes, which by definition means that for any node x in T_r , at most one of its children is a k -neighbour of x . Choose P to start at the root r , and as long as the last node in P has a k -neighbour v , v is appended to P . This set of nodes is obviously a path in T_r . No node in P can possibly have a k -neighbour outside of P , therefore all connected components of $T \setminus N[P]$ have LMIM-width $\leq k - 1$. Furthermore, all components of $T - N[P]$ are full rooted sub-trees of T_r and so the labels are already known.

Type 2 trees: In these trees the root r is k -critical. We look at the trees rooted in the two k -neighbours of r , $T_r[v_1]$ and $T_r[v_2]$. By Remark 2 these must both be Type 1 trees, and so we find paths P_1, P_2 in $T_r[v_1]$ and $T_r[v_2]$ respectively, as described above. Gluing these paths together at r we get a satisfying path for T_r , and we still have correct labels for the components $T \setminus N[P]$.

Type 3 trees: In these trees, r has exactly one child v such that $T_r[v]$ is of type 2 and none of its other children have LMIM-width k . We choose P as we did above for $T_r[v]$. r is clearly not a k -neighbour of v , or else $D_T(v, k) = 3$. Every other node in P has all their neighbours in $T_r[v]$. Again, every tree in $T \setminus N[P]$ is a full rooted subtree, and every label is known.

Type 4 trees: In these trees, T_r contains precisely one node $w \neq r$ such that w is the parent of a k -critical node, x . This w is easy to find using the labels, and clearly the tree $T_r[w]$ is a type 3 tree with LMIM-width k . We find a path P that is satisfying in $T_r[w]$ as described above. w is still not a k -neighbour of x , therefore P is a satisfying path. In this case, we have one connected component of $T \setminus N[P]$ that is not a full rooted subtree of T_r , that is $T_r \setminus T_r[w]$. Thus for every ancestor y of w (the blue path in Figure 6) $T_r[y] \setminus T_r[w]$ is not a full rooted subtree either, and we need to update the labels of these trees. However, $T_r[y] \setminus T_r[w]$ is by definition equal to $T_r[y, k]$, whose label is equal to $label(T_r[y])$ without its first number. Thus we quickly find the correct labels to do the recursive call.




Fig. 6. The path P in green for the proof of Theorem 3.

5 Conclusion

We have given an $O(n \log n)$ algorithm computing the LMIM-width and an optimal layout of an n -node tree. This is the first graph class of LMIM-width larger than 1 having a polynomial-time algorithm computing LMIM-width and thus constitutes an important step towards a better understanding of LMIM-width. Indeed, for the development of FPT algorithms computing tree-width and path-width of general graphs, one could argue that the algorithm of [6] computing optimal path-decompositions of a tree in time $O(n \log n)$ was a stepping stone. The situation is different for MIM-width and LMIM-width, as it is W-hard to compute these parameters [18], but it is similar in the sense that our objective has been to achieve an understanding of how to take a graph and assemble a decomposition of it, in this case a linear one, such that it has cuts of low MIM. To achieve this objective a polynomial-time algorithm for trees has been our main goal.

Bibliography

- [1] Adler, I. and Kanté, M. M. (2015). Linear rank-width and linear clique-width of trees. *Theor. Comput. Sci.*, 589:87–98.
- [2] Belmonte, R. and Vatshelle, M. (2013). Graph classes with structured neighborhoods and algorithmic applications. *Theor. Comput. Sci.*, 511:54 – 65.
- [3] Bergougnoux, B. and Kanté, M. M. (2018). Rank based approach on graphs with structured neighborhood. *CoRR*, abs/1805.11275.
- [4] Bui-Xuan, B.-M., Telle, J. A., and Vatshelle, M. (2013). Fast dynamic programming for locally checkable vertex subset and vertex partitioning problems. *Theor. Comput. Sci.*, 511:66 – 76.
- [5] Diestel, R. (2012). *Graph Theory, 4th Edition*, volume 173 of *Graduate texts in mathematics*. Springer.
- [6] Ellis, J. A., Sudborough, I. H., and Turner, J. S. (1994). The vertex separation and search number of a graph. *Inf. Comput.*, 113(1):50–79.
- [7] Fomin, F. V., Golovach, P. A., and Raymond, J.-F. (2018). On the tractability of optimization problems on H-graphs. In *Proc. ESA 2018*, pages 30:1 – 30:14.
- [8] Galby, E., Munaro, A., and Ries, B. (2018). Semitotal domination: New hardness results and a polynomial-time algorithm for graphs of bounded mim-width. *CoRR*, abs/1810.06872.
- [9] Golovach, P. A., Heggernes, P., Kanté, M. M., Kratsch, D., Sæther, S. H., and Villanger, Y. (2018). Output-polynomial enumeration on graphs of bounded (local) linear mim-width. *Algorithmica*, 80(2):714–741.
- [10] Golumbic, M. C. and Rotics, U. (2000). On the clique-width of some perfect graph classes. *Int. J. Found. Comput. Sci.*, 11(3):423–443.
- [11] Hliněný, P., Oum, S., Seese, D., and Gottlob, G. (2008). Width parameters beyond tree-width and their applications. *Comput. J.*, 51(3):326–362.
- [12] Jaffke, L., Kwon, O., Strømme, T. J. F., and Telle, J. A. (2018a). Generalized distance domination problems and their complexity on graphs of bounded mim-width. In *13th International Symposium on Parameterized and Exact Computation, IPEC 2018, August 20-24, 2018, Helsinki, Finland*, pages 6:1–6:14.
- [13] Jaffke, L., Kwon, O., and Telle, J. A. (2017). Polynomial-time algorithms for the longest induced path and induced disjoint paths problems on graphs of bounded mim-width. In *12th International Symposium on Parameterized and Exact Computation, IPEC 2017, September 6-8, 2017, Vienna, Austria*, pages 21:1–21:13.
- [14] Jaffke, L., Kwon, O., and Telle, J. A. (2018b). A unified polynomial-time algorithm for feedback vertex set on graphs of bounded mim-width. In *35th Symposium on Theoretical Aspects of Computer Science, STACS 2018, February 28 to March 3, 2018, Caen, France*, pages 42:1–42:14.
- [15] Mengel, S. (2018). Lower bounds on the mim-width of some graph classes. *Discrete Applied Mathematics*, 248:28–32.
- [16] Möhring, R. H. (1990). Graph problems related to gate matrix layout and pla folding. In *Computational graph theory*, pages 17–51. Springer.

- [17] Oum, S. (2017). Rank-width: Algorithmic and structural results. *Discrete Applied Mathematics*, 231:15–24.
- [18] Sæther, S. H. and Vatshelle, M. (2016). Hardness of computing width parameters based on branch decompositions over the vertex set. *Theor. Comput. Sci.*, 615:120–125.
- [19] Skodinis, K. (2003). Construction of linear tree-layouts which are optimal with respect to vertex separation in linear time. *J. Algorithms*, 47(1):40–59.
- [20] Vatshelle, M. (2012). *New width parameters of graphs*. PhD thesis, University of Bergen, Norway.
- [21] Yamazaki, K. (2018). Inapproximability of rank, clique, boolean, and maximum induced matching-widths under small set expansion hypothesis. *Algorithms*, 11(11):173.